

Concurrent Read Exclusive Write

Concurrent-read exclusive-write (CREW) is a fundamental concurrent programming paradigm. As the name suggests, it is about implementing a system that allows for multiple threads to read a shared object, but only allows one thread to write to a shared object at any given moment. In this part of the assignment, your task is to write a class called **WordCounter**, whose work is described below:

1. This class will read a list of plain text files, all of which are in a given folder.
2. Based on the word counts, it will generate a table that shows which word appears how many times and in which file.

Let us look at a worked out example. Suppose there are three text files:

- text-a.txt: "There are so many words in this file. Well, actually not that many if you really start counting them."
- text-b.txt: "The words are scattered everywhere, and in no order whatsoever. What. So. Ever. How will counting them help?"
- text-c.txt: "So, what can you do?"

The output *table* is shown below.

For the purpose of this assignment, you need not worry about hyphenated words (e.g., "e-mail"), apostrophes (e.g., "it's" or "Ma'am"), 'word-like' objects (e.g., a date written as "01-12-2020"), acronyms and abbreviations (e.g., "Mr.", "U.S.A."), or quotes.

This homework will not be tested with very large number of files or very large number of threads. You can safely assume that there will be at most 20 files and at most 20 threads that the grading process will use.

Note the following *required features* of the output:

- the words in column one are alphabetically ordered
- the columns corresponding to the files are alphabetically ordered by file name (the file names used for grading will all be lower-cased, so you don't have to worry about special characters)
- the first column's width is the longest word's length plus one character (so that there's a single space between the longest word and the first number)
- all the columns corresponding to files are separated by 4 spaces, and each column counts the number of occurrences of a word in *that* file
- the last column counts the total number of occurrences of a word across *all* files
- all the counts are left-aligned in the columns
- all words are lower-cased

- punctuation (. , ; ! ?) has been completely ignored

The **WordCounter** file must generate this output for the set of all files in a specific folder. To accomplish this, it must read each text file in a separate thread. Of course, all the words and word-counts from individual files have to be consolidated to write the final table. This is where you must carefully put together the results obtained by the different threads.

First, you should test this code without multithreading (i.e., the number of threads is just one -- the main thread) and then, with multithreading, to see what kind of speed boost (if any) you are getting. This will also help you identify bugs in your code. If something is going wrong with a single thread, then of course, the multithreaded version will also have bugs (and debugging single-threaded programs are much easier).

The basic structure of WordCounter is given below. Your class must have the constants **FOLDER_OF_TEXT_FILES**, **WORD_COUNT_TABLE_FILE**, and **NUMBER_OF_THREADS**. These are the only values we will change to run your code, so you have to be absolutely sure that your code is modular and robust enough to not crash if these details are changed when the code is run on another machine. For grading, only these constants will be changed, and your code will then be run (the graders will not modify anything else).

```
public class WordCounter {

    // The following are the ONLY variables we will modify for grading.
    // The rest of your code must run with no changes.
    public static final Path FOLDER_OF_TEXT_FILES = Paths.get("..."); // path to the folder
    where input text files are located
    public static final Path WORD_COUNT_TABLE_FILE = Paths.get("..."); // path to the output
    plain-text (.txt) file
    public static final int NUMBER_OF_THREADS    = 2;                // max. number of threads
    to spawn

    public static void main(String... args) {
        // your implementation of how to run the WordCounter as a stand-alone multi-threaded
        program
    }
}
```